

Load	5,348	17.90	0.00%	8.56 ms	30 ms	68 ms	178 ms
Stress	8,522	47.55	0.00%	25.06 ms	139 ms	339 ms	648 ms
Spike	3,711	18.77	0.00%	22.05 ms	97 ms	403 ms	622 ms
Endurance	55,622	92.89	0.00%	19.04 ms	83 ms	293 ms	1,031 ms

It reported backend CPU/RSS average and peak values as Load 1.99%/17.3% and 33,530/45,744 KB; Stress 9.17%/24.6% and 42,382/64,304 KB; Spike 3.72%/18.0% and 34,357/63,664 KB; Endurance 7.76%/26.9% and 57,294/76,000 KB.

The AI suggested CI gates of error rate at most 1%, endurance p95 at most 150 ms, p99 at most 600 ms, and throughput at least 85 RPS. It classified an index on `orders(user\_id, id)`, an index on `users(email)`, and SQLite WAL as feasible experiments. It classified a normal connection pool as unsuitable for this local SQLite SUT and said caching is not justified by the current evidence. It also stated that 92.89 RPS is a tested stable level, not a proven maximum hardware capacity.

My review and correction

I recalculated values from the JTL elapsed column. The Load, Stress, and Endurance percentile values agreed. The AI's merged Spike p95 was not exact: it wrote 97 ms, but the sorted `results/spike.jtl` value is ****93 ms****. It also wrote 105 ms for the spike-stage p95; the stage JTL value is ****103 ms****. I use the direct raw values in the main report. I use 92.70 RPS in the report because I divide by the planned 600 seconds; the AI divided by its observed elapsed window. This is a denominator difference, so neither value proves a capacity limit.

Human responsibility

I verified that the resource-monitor PID was `/opt/homebrew/bin/node server.js`. I kept the raw files, reported only measured values, and treated the coupon response as a functional bug only after a separate reproduction.